# Hadoop job_201104051408_0002 on master

**User:** root
**Job Name:** streamjob4686292789010298122.jar
**Job File:** hdfs://master:54310/mnt/hadoop/hadoop-mapred/mapred/staging/root/.staging/job_201104051408_0002/job.xml
**Submit Host:** master
**Submit Host Address:** 10.252.78.192
**Job-ACLs: All users are allowed**
**Job Setup:** Successful
**Status:** Succeeded
**Started at:** Tue Apr 05 15:16:58 UTC 2011
**Finished at:** Tue Apr 05 15:32:42 UTC 2011
**Finished in:** 15mins, 44sec
**Job Cleanup:** Successful

| Kind | % Complete | Num Tasks | Pending | Running | Complete | Killed | Failed/Killed Task Attempts |
|---|---|---|---|---|---|---|---|
| map | 100.00% | 60 | 0 | 0 | 60 | 0 | 0 / 3 |
| reduce | 100.00% | 20 | 0 | 0 | 20 | 0 | 0 / 3 |

| | Counter | Map | Reduce | Total |
|---|---|---|---|---|
| Job Counters | SLOTS_MILLIS_MAPS | 0 | 0 | 3,987,219 |
| | Launched reduce tasks | 0 | 0 | 23 |
| | Total time spent by all reduces waiting after reserving slots (ms) | 0 | 0 | 0 |
| | Total time spent by all maps waiting after reserving slots (ms) | 0 | 0 | 0 |
| | Launched map tasks | 0 | 0 | 63 |
| | Data-local map tasks | 0 | 0 | 63 |
| | SLOTS_MILLIS_REDUCES | 0 | 0 | 7,719,500 |
| FileSystemCounters | FILE_BYTES_READ | 13,108,747,874 | 10,158,116,315 | 23,266,864,189 |
| | HDFS_BYTES_READ | 3,574,093,330 | 0 | 3,574,093,330 |
| | FILE_BYTES_WRITTEN | 23,240,527,552 | 10,159,130,245 | 33,399,657,797 |
| | HDFS_BYTES_WRITTEN | 0 | 3,624,808,508 | 3,624,808,508 |
| Map-Reduce Framework | Reduce input groups | 0 | 29,695,820 | 29,695,820 |
| | Combine output records | 0 | 0 | 0 |
| | Map input records | 9,113,764 | 0 | 9,113,764 |

| Reduce shuffle bytes | 0 | 10,121,914,872 | 10,121,914,872 |
|---|---|---|---|
| Reduce output records | 0 | 9,113,764 | 9,113,764 |
| Spilled Records | 408,833,268 | 178,202,294 | 587,035,562 |
| Map output bytes | 9,795,495,383 | 0 | 9,795,495,383 |
| Map input bytes | 3,573,386,274 | 0 | 3,573,386,274 |
| SPLIT_RAW_BYTES | 6,600 | 0 | 6,600 |
| Map output records | 178,202,294 | 0 | 178,202,294 |
| Combine input records | 0 | 0 | 0 |
| Reduce input records | 0 | 178,202,294 | 178,202,294 |

Map Completion Graph - close

Reduce Completion Graph - close

Go back to JobTracker

Cloudera's Distribution for Hadoop, 2011.