

# Hadoop job\_201104051408\_0005 on master

**User:** root

**Job Name:** streamjob6417307444530919499.jar

**Job File:** [hdfs://master:54310/mnt/hadoop/hadoop-mapred/mapred/staging/root/.staging/job\\_201104051408\\_0005/job.xml](hdfs://master:54310/mnt/hadoop/hadoop-mapred/mapred/staging/root/.staging/job_201104051408_0005/job.xml)

**Submit Host:** master

**Submit Host Address:** 10.252.78.192

**Job-ACLs:** All users are allowed

**Job Setup:** Successful

**Status:** Succeeded

**Started at:** Tue Apr 05 16:04:41 UTC 2011

**Finished at:** Tue Apr 05 16:19:22 UTC 2011

**Finished in:** 14mins, 40sec


**Job Cleanup:** Successful

Kind	% Complete	Num Tasks	Pending	Running	Complete	Killed	Failed/Killed Task Attempts
map	100.00%	60	0	0	60	0	0 / 3
reduce	100.00%	20	0	0	20	0	0 / 1


	Counter	Map	Reduce	Total
Job Counters	SLOTS_MILLIS_MAPS	0	0	3,724,463
	Launched reduce tasks	0	0	21
	Total time spent by all reduces waiting after reserving slots (ms)	0	0	0
	Rack-local map tasks	0	0	3
	Total time spent by all maps waiting after reserving slots (ms)	0	0	0
	Launched map tasks	0	0	63
	Data-local map tasks	0	0	60
	SLOTS_MILLIS_REDUCE	0	0	7,339,052
FileSystemCounters	FILE_BYTES_READ	12,970,498,308	10,311,830,018	23,282,328,326
	HDFS_BYTES_READ	3,626,440,528	0	3,626,440,528
	FILE_BYTES_WRITTEN	23,256,281,350	10,312,843,948	33,569,125,298
	HDFS_BYTES_WRITTEN	0	3,626,239,443	3,626,239,443
Map-Reduce Framework	Reduce input groups	0	29,695,820	29,695,820
	Combine output records	0	0	0

Map input records	9,113,764	0	9,113,764
Reduce shuffle bytes	0	10,247,584,725	10,247,584,725
Reduce output records	0	9,113,764	9,113,764
Spilled Records	403,590,388	178,202,294	581,792,682
Map output bytes	9,949,209,092	0	9,949,209,092
Map input bytes	3,626,261,856	0	3,626,261,856
SPLIT_RAW_BYTES	6,600	0	6,600
Map output records	178,202,294	0	178,202,294
Combine input records	0	0	0
Reduce input records	0	178,202,294	178,202,294

Map Completion Graph - [close](#)



Reduce Completion Graph - [close](#)



[Go back to JobTracker](#)

---

[Cloudera's Distribution for Hadoop](#), 2011.