# Hadoop job_201104052005_0001 on master

**User:** root
**Job Name:** streamjob7330947814582378055.jar
**Job File:** hdfs://master:54310/mnt/hadoop/hadoop-mapred/mapred/staging/root/.staging/job_201104052005_0001/job.xml
**Submit Host:** master
**Submit Host Address:** 10.124.175.97
**Job-ACLs: All users are allowed**
**Job Setup:** Successful
**Status:** Succeeded
**Started at:** Tue Apr 05 20:24:54 UTC 2011
**Finished at:** Tue Apr 05 20:28:44 UTC 2011
**Finished in:** 3mins, 50sec
**Job Cleanup:** Successful

| Kind | % Complete | Num Tasks | Pending | Running | Complete | Killed | Failed/Killed Task Attempts |
|------|-----------|-----------|---------|---------|----------|--------|------------------------------|
| map | 100.00% | 317 | 0 | 0 | 317 | 0 | 0 / 21 |
| reduce | 100.00% | 100 | 0 | 0 | 100 | 0 | 0 / 22 |

| | Counter | Map | Reduce | Total |
|---|---------|-----|--------|-------|
| Job Counters | SLOTS_MILLIS_MAPS | 0 | 0 | 4,913,414 |
| | Launched reduce tasks | 0 | 0 | 122 |
| | Total time spent by all reduces waiting after reserving slots (ms) | 0 | 0 | 0 |
| | Rack-local map tasks | 0 | 0 | 14 |
| | Total time spent by all maps waiting after reserving slots (ms) | 0 | 0 | 0 |
| | Launched map tasks | 0 | 0 | 338 |
| | Data-local map tasks | 0 | 0 | 324 |
| | SLOTS_MILLIS_REDUCES | 0 | 0 | 7,495,078 |
| FileSystemCounters | FILE_BYTES_READ | 0 | 3,558,351,099 | 3,558,351,099 |
| | HDFS_BYTES_READ | 21,161,567,888 | 0 | 21,161,567,888 |
| | FILE_BYTES_WRITTEN | 3,575,399,600 | 3,563,419,889 | 7,138,819,489 |
| | HDFS_BYTES_WRITTEN | 0 | 3,573,386,274 | 3,573,386,274 |
| Map-Reduce Framework | Reduce input groups | 0 | 9,113,764 | 9,113,764 |
| | Combine output records | 0 | 0 | 0 |

| Map input records | 9,113,764 | 0 | 9,113,764 |
|---|---|---|---|
| Reduce shuffle bytes | 0 | 3,554,613,703 | 3,554,613,703 |
| Reduce output records | 0 | 9,113,764 | 9,113,764 |
| Spilled Records | 9,113,764 | 9,113,764 | 18,227,528 |
| Map output bytes | 3,533,960,482 | 0 | 3,533,960,482 |
| Map input bytes | 21,154,547,097 | 0 | 21,154,547,097 |
| SPLIT_RAW_BYTES | 40,893 | 0 | 40,893 |
| Map output records | 9,113,764 | 0 | 9,113,764 |
| Combine input records | 0 | 0 | 0 |
| Reduce input records | 0 | 9,113,764 | 9,113,764 |

Map Completion Graph - close
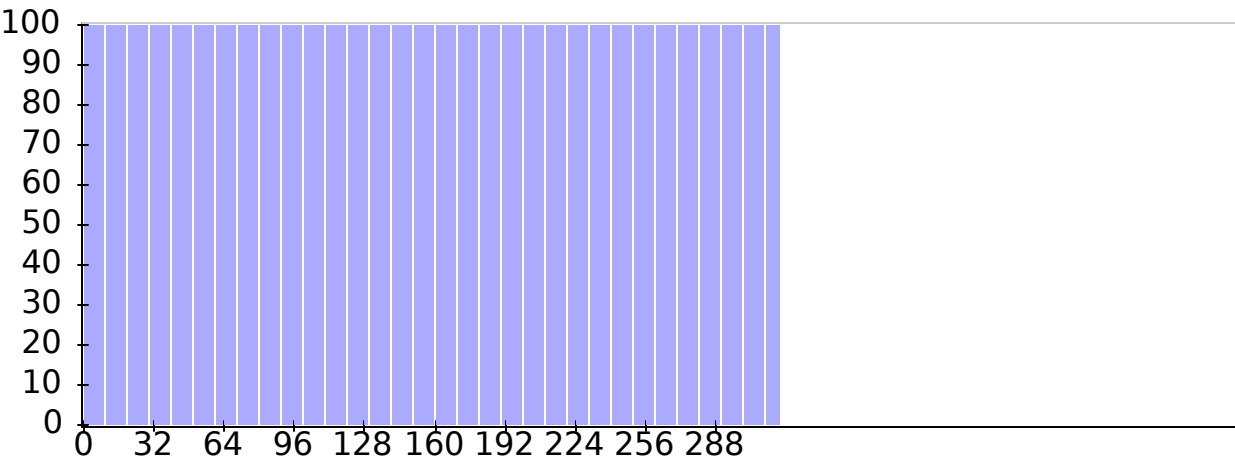


Reduce Completion Graph - close

Go back to JobTracker

Cloudera's Distribution for Hadoop, 2011.