

# Hadoop job\_201104052005\_0002 on master

**User:** root

**Job Name:** streamjob511062283573883239.jar

**Job File:** [hdfs://master:54310/mnt/hadoop/hadoop-mapred/mapred/staging/root/.staging/job\\_201104052005\\_0002/job.xml](hdfs://master:54310/mnt/hadoop/hadoop-mapred/mapred/staging/root/.staging/job_201104052005_0002/job.xml)

**Submit Host:** master

**Submit Host Address:** 10.124.175.97

**Job-ACLs:** All users are allowed

**Job Setup:** Successful

**Status:** Succeeded

**Started at:** Tue Apr 05 20:28:48 UTC 2011

**Finished at:** Tue Apr 05 20:32:40 UTC 2011

**Finished in:** 3mins, 52sec


**Job Cleanup:** Successful

Kind	% Complete	Num Tasks	Pending	Running	Complete	Killed	Failed/Killed Task Attempts
map	100.00%	100	0	0	100	0	0 / 13
reduce	100.00%	100	0	0	100	0	0 / 16


	Counter	Map	Reduce	Total
Job Counters	SLOTS_MILLIS_MAPS	0	0	3,667,948
	Launched reduce tasks	0	0	116
	Total time spent by all reduces waiting after reserving slots (ms)	0	0	0
	Rack-local map tasks	0	0	9
	Total time spent by all maps waiting after reserving slots (ms)	0	0	0
	Launched map tasks	0	0	113
	Data-local map tasks	0	0	104
	SLOTS_MILLIS_REDUCE	0	0	6,754,608
FileSystemCounters	FILE_BYTES_READ	10,302,520,337	10,158,116,663	20,460,637,000
	HDFS_BYTES_READ	3,573,397,274	0	3,573,397,274
	FILE_BYTES_WRITTEN	20,322,032,016	10,163,186,553	30,485,218,569
	HDFS_BYTES_WRITTEN	0	3,624,808,507	3,624,808,507
Map-Reduce Framework	Reduce input groups	0	29,695,820	29,695,820
	Combine output records	0	0	0

Map input records	9,113,764	0	9,113,764
Reduce shuffle bytes	0	10,026,784,745	10,026,784,745
Reduce output records	0	9,113,764	9,113,764
Spilled Records	356,404,588	178,202,294	534,606,882
Map output bytes	9,795,495,383	0	9,795,495,383
Map input bytes	3,573,386,274	0	3,573,386,274
SPLIT_RAW_BYTES	11,000	0	11,000
Map output records	178,202,294	0	178,202,294
Combine input records	0	0	0
Reduce input records	0	178,202,294	178,202,294

Map Completion Graph - [close](#)



Reduce Completion Graph - [close](#)



[Go back to JobTracker](#)

---

[Cloudera's Distribution for Hadoop](#), 2011.